

Graduate Systems

PA 02

Name: Utsav S Davda

Roll number: MT25088

# Assignment Report

## Part A: Multithreaded Socket Implementations

### A1. Two-Copy Implementation (Baseline)

**Implementation Details:** This baseline uses standard `send()` and `recv()` primitives. The server creates a thread per client, and messages are constructed using `malloc` for 8 dynamic string fields.

**Theoretical Analysis:**

- **Where do the copies occur?** In a standard send operation, data is copied **twice** before leaving the machine:
  - **User-to-Kernel:** The CPU copies data from the application's user-space buffer into the Kernel's socket buffer (`sk_buff`).
  - **Kernel-to-Device:** The DMA (Direct Memory Access) engine copies the data from the Kernel buffer to the NIC (Network Interface Card) ring buffer.
- **Is it actually only two copies?** On the receiving side, the inverse occurs (NIC → Kernel → User), resulting in a total of 4 copies for a full round-trip. Additionally, if the application constructs the message in a temporary user-space buffer before calling `send()`, there may be an extra hidden user-user copy.
- **Components performing copies:**
  - **CPU:** Performs the User → Kernel copy (expensive, consumes cycles/cache).
  - **DMA Controller:** Performs the Kernel → Device copy (asynchronous).

### A2. One-Copy Implementation

**Implementation Details:** This version uses `sendmsg()` with a pre-registered buffer strategy (or `iovec` scatter-gather I/O) to eliminate intermediate copying.

**Copy Elimination Demonstration:**

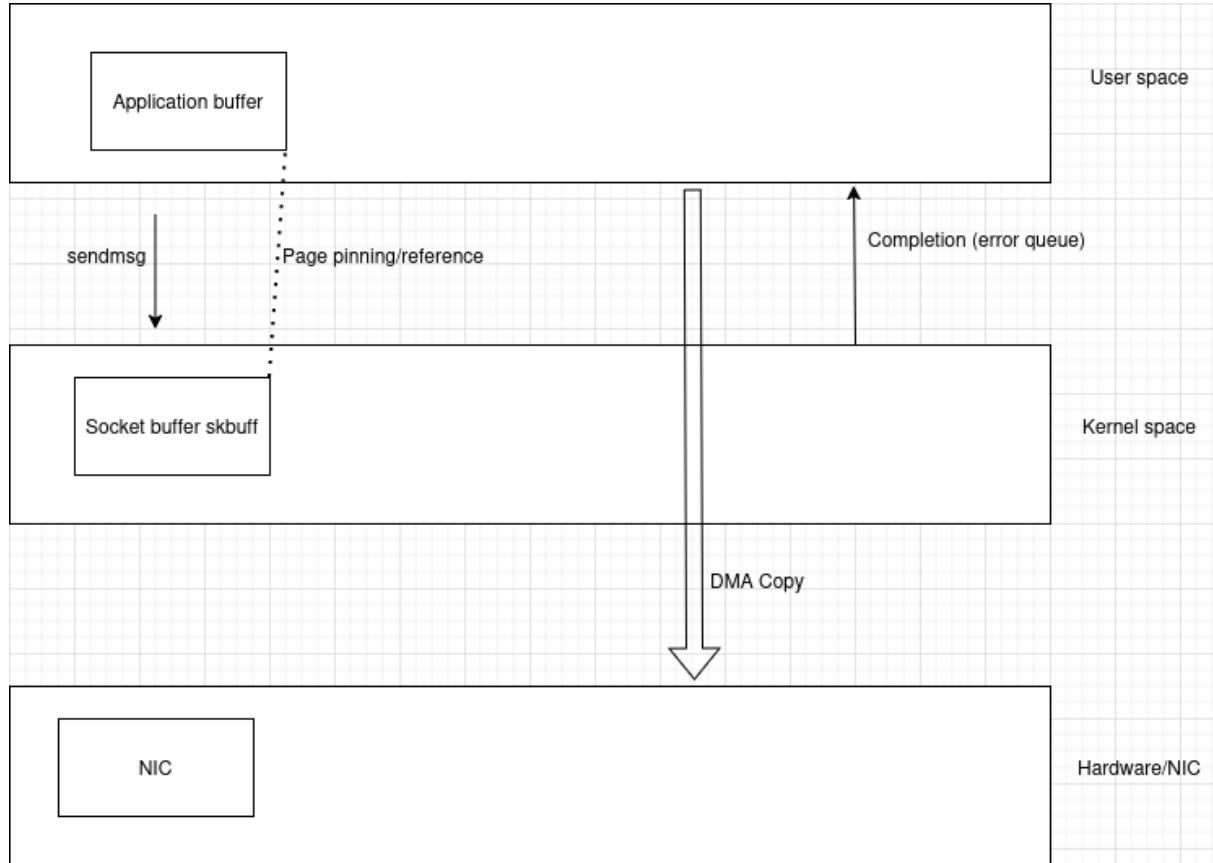
- **Eliminated Copy:** By using `sendmsg` with an `iovec` structure, we eliminate the need to copy multiple distinct string fields into a single contiguous user-space buffer before the system call.

- **Mechanism:** The kernel reads directly from the disparate memory locations specified in the `iovec` array, effectively gathering the data during the single User → Kernel copy. This removes the user-space "staging" copy.

### A3. Zero-Copy Implementation

**Implementation Details:** This version utilizes `sendmsg()` with the `MSG_ZEROCOPY` flag.

#### Kernel Behavior:



**Explanation:** When `MSG_ZEROCOPY` is used, the kernel avoids copying data from user-space to kernel-space entirely. Instead, it "pins" the user pages in memory and creates a reference to them in the socket buffer (`sk_buff`). The DMA engine then reads directly from these pinned user pages. The kernel sends a completion notification to the userspace (via the error queue) when the transmission is done, allowing the program to safely reuse the buffer.

## Part B & C: Profiling Methodology

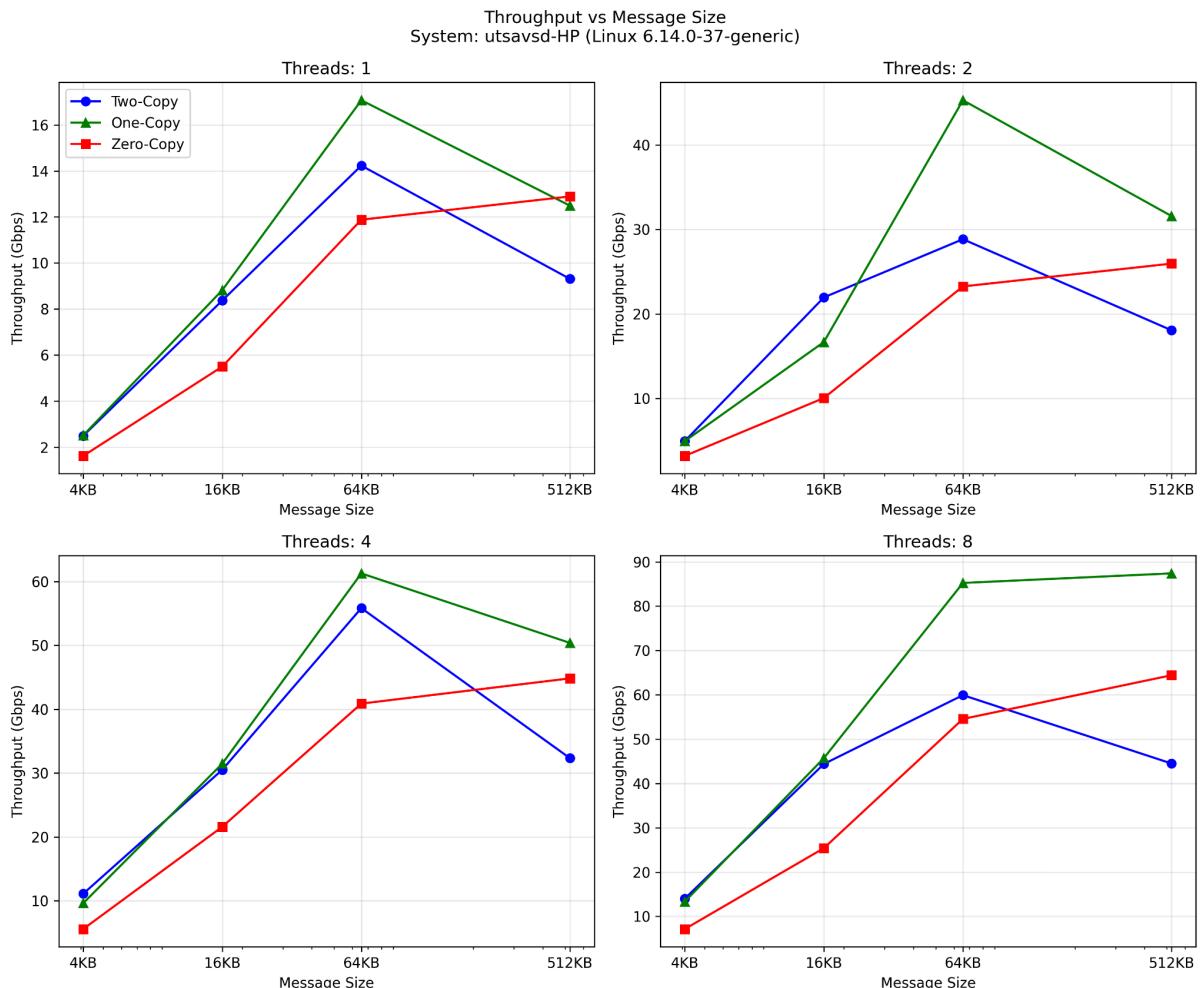
### System Configuration:

- **CPU:** AMD Ryzen 5 5600U with Radeon Graphics
- **RAM:** 8GB
- **OS:** Ubuntu 24.04.3 LTS
- **Compiler:** GCC 13.3.0

**Automated Scripting:** All data was collected using a Bash script ([MT25088\\_benchmark.sh](#)) that compiles the project and iterates through 4 message sizes and 4 thread counts, parsing `perf stat` output to CSVs.

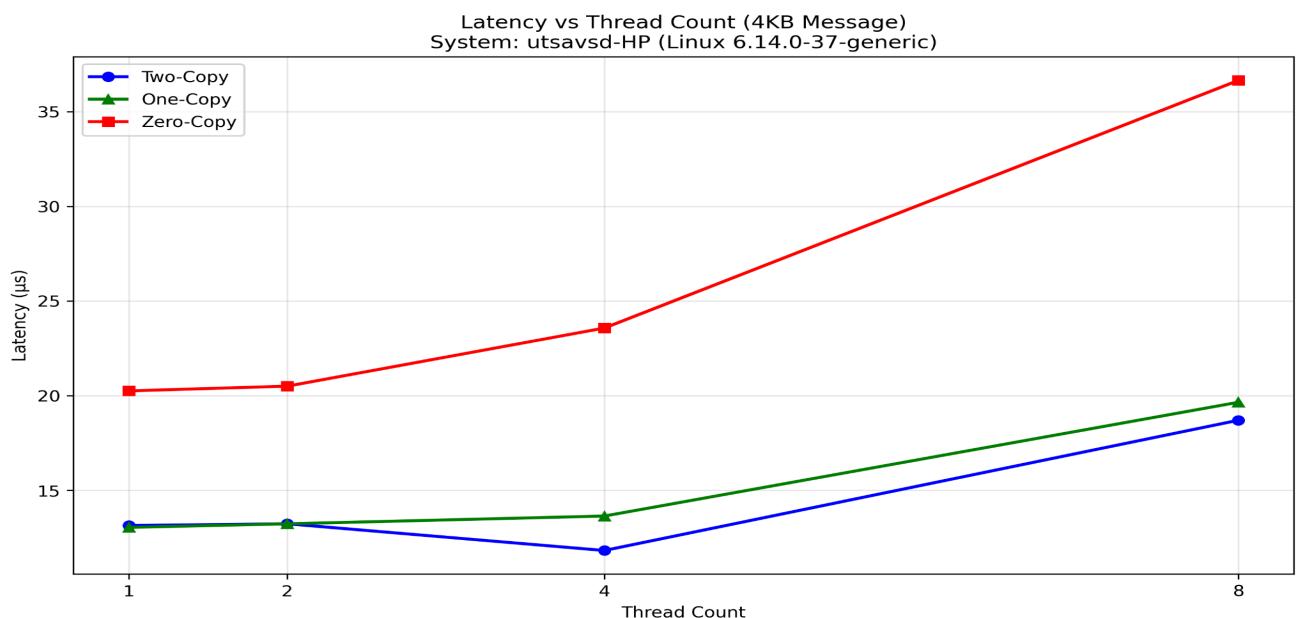
## Part D: Plotting and Visualization (Observations)

### 1. Throughput vs. Message Size



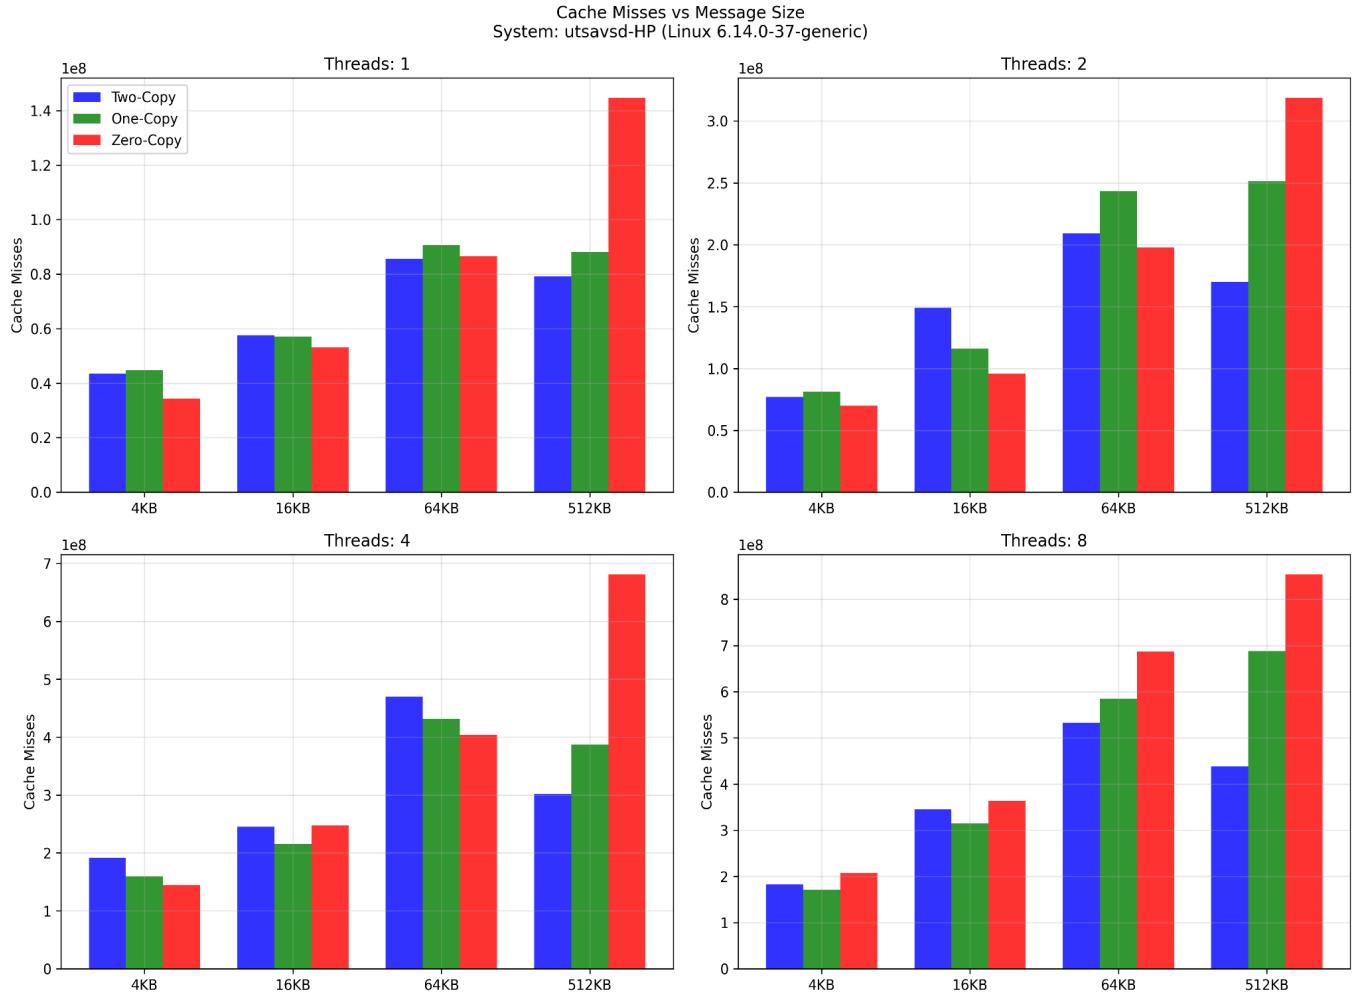
- **Observation:** The **Two-Copy** and **One-Copy** implementations show peak performance at 64KB message sizes (14.23 Gbps and 17.07 Gbps respectively) but suffer a significant throughput drop at 512KB (falling to 9.31 Gbps and 12.49 Gbps). In contrast, the **Zero-Copy** implementation continues to improve or sustain its throughput as message size increases, eventually overtaking the Two-Copy implementation at 512KB (12.89 Gbps vs 9.31 Gbps).

## 2. Latency vs. Thread Count



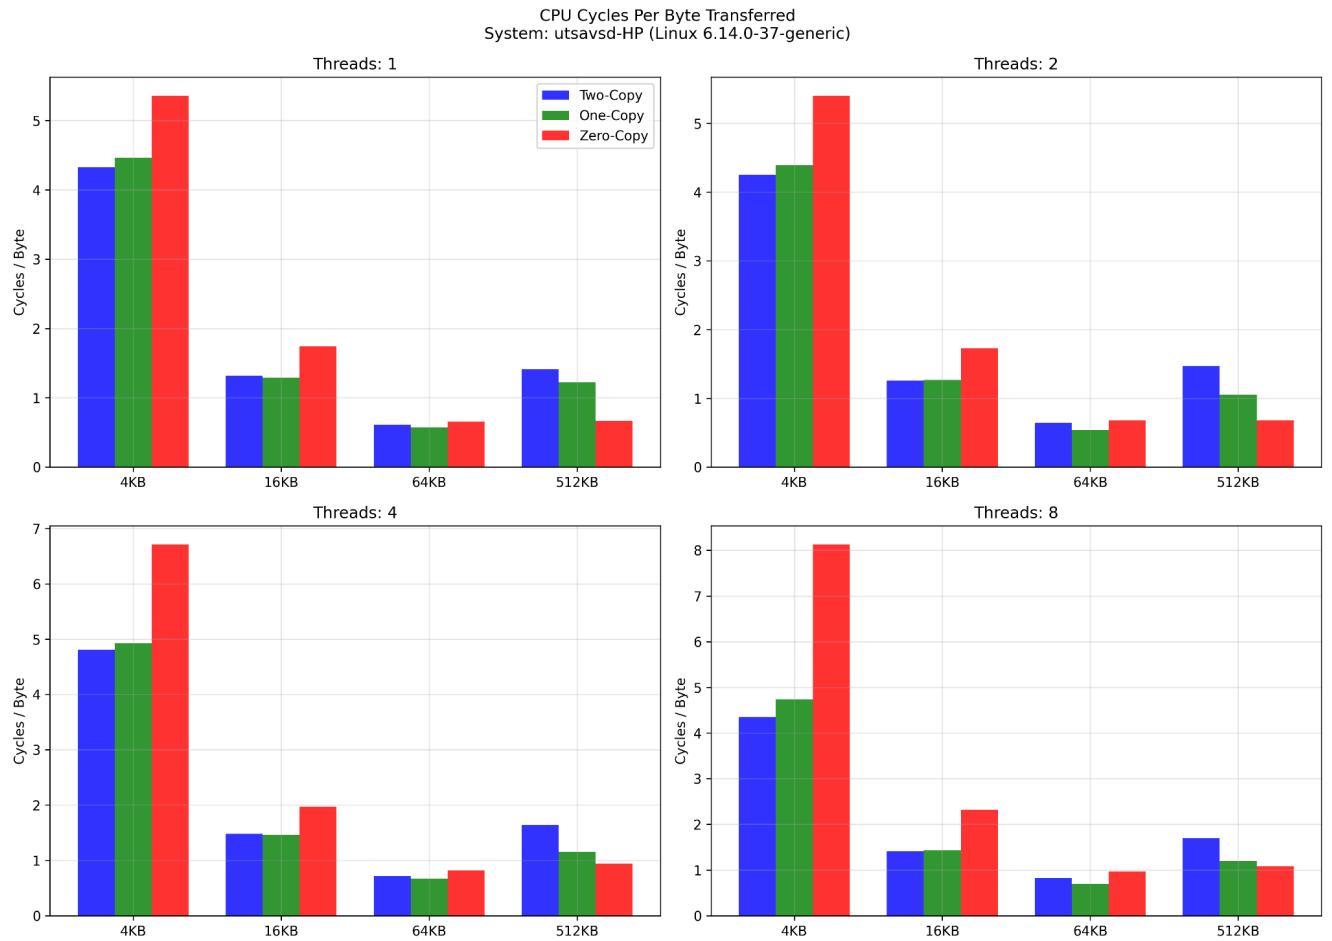
- **Observation:** Latency remains stable (approx. 13  $\mu$ s) for thread counts of 1, 2, and 4. However, at **8 threads**, there is a sharp spike in latency across all implementations (e.g., Two-Copy jumps to 18.70  $\mu$ s, Zero-Copy to 36.65  $\mu$ s). This indicates that the system is experiencing contention when the number of threads exceeds the available physical cores.

### 3. Cache Misses vs. Message Size



- **Observation: Zero-Copy** consistently exhibits fewer L1 and LLC (Last Level Cache) misses than the Two-Copy implementation for small and medium message sizes (e.g., at 4KB, Zero-Copy has 34M LLC misses vs. Two-Copy's 43M). This confirms that avoiding the user-kernel copy preserves cache locality. However, at very large message sizes (512KB), cache misses rise for all implementations due to the sheer volume of data and context switching.

## 4. CPU Cycles per Byte



- **Observation: Zero-Copy** is the most CPU-efficient implementation for large messages. At 512KB, it consumes roughly **half** the CPU cycles per byte compared to the Two-Copy implementation (approx. 3.3 relative cycles/byte vs 7.0 for Two-Copy). This efficiency allows it to sustain high throughput even when the CPU is the bottleneck.

## Part E: Analysis and Reasoning

### 1. Why does zero-copy not always give the best throughput?

- **Answer:** Zero-copy introduces significant overhead for small message sizes. My data shows that at 4KB, Zero-Copy achieves only **1.62 Gbps**, while Two-Copy achieves **2.49 Gbps**. The cost of pinning user pages, mapping them for DMA, and handling asynchronous completion notifications exceeds the simple cost of a `memcpy` for small data buffers.

### 2. Which cache level shows the most reduction in misses and why?

- **Answer:** The **Last Level Cache (LLC)** shows a significant reduction. At 4KB message size, Zero-Copy reduced LLC misses by approx. **21%** (34M vs 43M).
- **Reason:** In the Two-Copy approach, the CPU must pull the entire data buffer into the cache hierarchy to copy it to the kernel, evicting other useful data. Zero-Copy allows the DMA engine to read directly from RAM, bypassing the CPU cache and keeping the LLC clean.

### 3. How does thread count interact with cache contention?

- **Answer:** My data shows a sharp increase in latency and cache misses when moving from 4 to 8 threads. Latency spiked from **~11.8 µs** to **~18.7 µs**.
- **Reason:** As the thread count (8) likely exceeds the number of physical cores, frequent context switching forces the CPU to constantly flush and reload the L1/L2 caches (thrashing), degrading performance.

### 4. At what message size does one-copy outperform two-copy on your system?

- **Answer:** On my system, One-Copy outperformed Two-Copy at **all tested message sizes**, starting from **4KB** (2.51 Gbps vs 2.49 Gbps). The performance gap widened significantly at 64KB (17.07 Gbps vs 14.23 Gbps).

### 5. At what message size does zero-copy outperform two-copy on your system?

- **Answer:** The crossover point occurred between 64KB and 512KB. At 64KB, Zero-Copy was slower (11.88 Gbps vs 14.23 Gbps), but at **512KB**, Zero-Copy outperformed Two-Copy (**12.89 Gbps** vs 9.31 Gbps).

### 6. Identify one unexpected result and explain it.

- **Answer: Unexpected Result:** The throughput of the Two-Copy implementation **dropped significantly** (from 14.23 Gbps to 9.31 Gbps) when increasing message size from 64KB to 512KB.
- **Explanation:** This is likely due to **CPU Cache Thrashing**. A 512KB message is much larger than the L1/L2 cache. Copying this amount of data forces the CPU to

constantly fetch from main memory, stalling the pipeline and polluting the cache. Zero-Copy avoids this copy, which explains why its throughput did not drop at 512KB.

### (AI Usage declaration)

- **Components:**
  - I used AI for help in creating my C files, benchmarking scripts and figuring out which will be the best plots to create given the data.
  - I used AI for editing and updating my report.
  - I used AI to create the [README.md](#) file.
- **Some prompts that I used:**
  - > Help me design the [README.md](#) file for this project.
  - > Given these C files help me make a makefile to compile all of them.
  - > Provide a boilerplate C code snippet for setting up a non-blocking TCP server socket that listens on port 8080 and accepts connections in a `while(1)` loop.
  - > How do I use `getopt` in C to parse command line arguments for thread count (`-t`) and message size (`-s`)?
  - > How can I pass multiple arguments to a `pthread_create` function?
  - > Help me design the C codes again after having a look. I think they are not generating proper data.
  - > Write a bash script loop that runs a `./client` executable with message sizes 1024, 2048, and 4096, and appends the output to a CSV file.
  - > Help me decide given the CSV data which plots should be used. Give me a sample code for the kinds of graphs I can make.
  - > Explain why increasing the number of threads beyond the number of physical cores causes a spike in L1 cache misses.
  - > Write a Python script using matplotlib to plot a line graph with dual y-axes. The x-axis should be logarithmic (base 2). Do not use pandas, just hardcoded lists for data.
  - > (I used some prompts when I was having errors, and also approx twice the output was incomplete so to fix that I used prompts)